

① Given,

$$\underline{\underline{A}} = \begin{bmatrix} 5 & 2 \\ 2 & 2 \end{bmatrix}$$

a. Determining the population principal components  $\gamma_1$  and  $\gamma_2$  for the given covariance matrix:

Let  $I_2$  identity matrix be  $\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$

$$\therefore \lambda I = \begin{bmatrix} \lambda & 0 \\ 0 & \lambda \end{bmatrix}$$

Then,  $\det(A) - \det(I) = 0$

$$\text{or, } \det \begin{bmatrix} 5 & 2 \\ 2 & 2 \end{bmatrix} - \det \begin{bmatrix} \lambda & 0 \\ 0 & \lambda \end{bmatrix} = 0$$

$$\text{or, } \begin{vmatrix} 5-\lambda & 2-0 \\ 2-0 & 2-\lambda \end{vmatrix} = 0$$

$$\text{or, } (5-\lambda)(2-\lambda) - 2 \times 2 = 0 \quad (\text{def. of } ||)$$

$$\text{or, } \lambda^2 - 7\lambda + 6 = 0$$

$$\text{or, } \lambda^2 - 6\lambda - 1\lambda + 6 = 0$$

$$\text{or, } \lambda(\lambda-6) - 1(\lambda-6) = 0$$

either,  $\boxed{\lambda = 6}$  or  $\boxed{\lambda = 1}$



## Computing the Eigenvalues

If  $\lambda = 0$

$$Av = \lambda v$$

$$\therefore v(A - \lambda) = 0$$

$$\text{or, } \begin{bmatrix} -1 & 2 \\ 2 & -4 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\text{or, } \begin{bmatrix} -v_1 + 2v_2 \\ 2v_1 - 4v_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\therefore v = 2v_2$$

$$\text{or } \boxed{v_1 = 2v_2}$$

$$\therefore \text{Eigenvector is } \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$
$$u = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

### Normalizing

$$\hat{u} = \frac{u}{\|u\|} = \frac{1}{\sqrt{2^2 + 1^2}} = \frac{1}{\sqrt{5}}$$

$$\therefore u = \begin{bmatrix} 2/\sqrt{5} \\ 1/\sqrt{5} \end{bmatrix}$$

1<sup>st</sup> principal component

$$y = \frac{2}{\sqrt{5}}x_1 + \frac{1}{\sqrt{5}}x_2$$

2<sup>nd</sup> pc

$$y = -\frac{1}{\sqrt{5}}x_1 + \frac{2}{\sqrt{5}}x_2$$

If  $\lambda = 1$

$$Av = \lambda v$$

$$\therefore \begin{bmatrix} 4 & 2 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} v_1 \\ v_2 \end{bmatrix} = 0$$

$$\text{or, } \begin{bmatrix} 4v_1 + 2v_2 \\ 2v_1 + v_2 \end{bmatrix} = 0$$

$$\text{or, } 4v_1 = -2v_2$$

$$\text{or, } \boxed{v_2 = -2v_1}$$

$$\text{Eigenvector is } \begin{bmatrix} -1 \\ 2 \end{bmatrix}$$
$$w = \begin{bmatrix} -1 \\ 2 \end{bmatrix}$$

### Normalizing

$$\hat{w} = \frac{w}{\|w\|} = \frac{w}{\sqrt{1+4}}$$

$$\therefore \hat{w} = \begin{bmatrix} -1/\sqrt{5} \\ 2/\sqrt{5} \end{bmatrix}$$



b. Calculating the prop<sup>n</sup> of total pop<sup>n</sup>  
variance explained by 1<sup>st</sup> PC

$$\frac{\lambda_1}{\lambda_1 + \lambda_2} = \frac{6}{61} = 0.857$$

$$= 85.7\%$$

## NB\_hw\_2

October 9, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn import decomposition
from sklearn import preprocessing
from sklearn import metrics

%matplotlib inline
plt.style.use('seaborn-white')
```

```
[2]: # Load Data and Inspect
hand_digits = pd.read_csv('optdigits.tra', header=None)
print(hand_digits.info())
hand_digits.head(3)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3823 entries, 0 to 3822
Data columns (total 65 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    0      3823 non-null    int64
 1    1      3823 non-null    int64
 2    2      3823 non-null    int64
 3    3      3823 non-null    int64
 4    4      3823 non-null    int64
 5    5      3823 non-null    int64
 6    6      3823 non-null    int64
 7    7      3823 non-null    int64
 8    8      3823 non-null    int64
 9    9      3823 non-null    int64
10   10     3823 non-null    int64
11   11     3823 non-null    int64
12   12     3823 non-null    int64
13   13     3823 non-null    int64
14   14     3823 non-null    int64
15   15     3823 non-null    int64
```

16	16	3823 non-null	int64
17	17	3823 non-null	int64
18	18	3823 non-null	int64
19	19	3823 non-null	int64
20	20	3823 non-null	int64
21	21	3823 non-null	int64
22	22	3823 non-null	int64
23	23	3823 non-null	int64
24	24	3823 non-null	int64
25	25	3823 non-null	int64
26	26	3823 non-null	int64
27	27	3823 non-null	int64
28	28	3823 non-null	int64
29	29	3823 non-null	int64
30	30	3823 non-null	int64
31	31	3823 non-null	int64
32	32	3823 non-null	int64
33	33	3823 non-null	int64
34	34	3823 non-null	int64
35	35	3823 non-null	int64
36	36	3823 non-null	int64
37	37	3823 non-null	int64
38	38	3823 non-null	int64
39	39	3823 non-null	int64
40	40	3823 non-null	int64
41	41	3823 non-null	int64
42	42	3823 non-null	int64
43	43	3823 non-null	int64
44	44	3823 non-null	int64
45	45	3823 non-null	int64
46	46	3823 non-null	int64
47	47	3823 non-null	int64
48	48	3823 non-null	int64
49	49	3823 non-null	int64
50	50	3823 non-null	int64
51	51	3823 non-null	int64
52	52	3823 non-null	int64
53	53	3823 non-null	int64
54	54	3823 non-null	int64
55	55	3823 non-null	int64
56	56	3823 non-null	int64
57	57	3823 non-null	int64
58	58	3823 non-null	int64
59	59	3823 non-null	int64
60	60	3823 non-null	int64
61	61	3823 non-null	int64
62	62	3823 non-null	int64
63	63	3823 non-null	int64

```

64 64      3823 non-null   int64
dtypes: int64(65)
memory usage: 1.9 MB
None

```

```

[2]:
   0  1  2  3  4  5  6  7  8  9  ...  55  56  57  58  59  60  61  \
0  0  1  6 15 12  1  0  0  0  7  ...  0  0  0  6 14  7  1
1  0  0 10 16  6  0  0  0  0  7  ...  0  0  0 10 16 15  3
2  0  0  8 15 16 13  0  0  0  1  ...  0  0  0  9 14  0  0

```

```

62 63 64
0  0  0  0
1  0  0  0
2  0  0  7

```

[3 rows x 65 columns]

```

[3]: # reading the unique target variables
unique_targets = hand_digits[64].unique()
print(unique_targets)

```

```
[0 7 4 6 2 5 8 1 9 3]
```

The hand\_digits df has sixty-four ( $p = 64$ ) inputs plus the target variable that indicates the digit 0-9 as shown above.

```

[4]: #b Remove any unary column (i.e., containing only one value).
hand_digits = hand_digits.loc[:, hand_digits.nunique() > 1]
print(hand_digits.shape)

```

```
(3823, 63)
```

We dropped one column which was unary column.

```

[5]: # c Check for missing values
print(hand_digits.isnull().sum().sum())

```

```
0
```

There are not any missing values in the dataframe.

```

[6]: #d Checking if the data is standardized
if np.allclose(hand_digits.mean(), 0) and np.allclose(hand_digits.std(), 1):
    print("Data is standardized")
else:
    print("Data is not standardized")

```

```
Data is not standardized
```

```
[7]: # Standardize the data matrix X so that all variables are given a mean of zero
      ↪ and a standard deviation of one.
      # excluding the target variable (last column)
      from sklearn.preprocessing import StandardScaler
      X = hand_digits.iloc[:, :-1]
      scaler = StandardScaler()
      X = scaler.fit_transform(X)
```

```
[8]: standarized_df = pd.DataFrame(X)
```

```
[9]: # checking the oveall mean
      standarized_df.describe()
```

```
[9]:
```

	0	1	2	3	4 \
count	3.823000e+03	3.823000e+03	3.823000e+03	3.823000e+03	3.823000e+03
mean	6.870140e-16	-1.158721e-17	-1.623371e-16	2.239032e-16	-1.388142e-17
std	1.000131e+00	1.000131e+00	1.000131e+00	1.000131e+00	1.000131e+00
min	-3.476105e-01	-1.183724e+00	-2.771826e+00	-2.524041e+00	-9.809412e-01
25%	-3.476105e-01	-9.677879e-01	-4.239971e-01	-5.403346e-01	-9.809412e-01
50%	-3.476105e-01	-1.040426e-01	2.803515e-01	3.413125e-01	-2.682243e-01
75%	-3.476105e-01	7.597028e-01	7.499173e-01	7.821361e-01	8.008511e-01
max	8.880965e+00	2.271257e+00	9.847002e-01	1.002548e+00	1.869927e+00

	5	6	7	8	9 \
count	3.823000e+03	3.823000e+03	3.823000e+03	3.823000e+03	3.823000e+03
mean	-3.130833e-16	1.590265e-16	-2.157809e-16	-1.879509e-16	-3.210731e-16
std	1.000131e+00	1.000131e+00	1.000131e+00	1.000131e+00	1.000131e+00
min	-4.115666e-01	-1.353324e-01	-2.362917e-02	-6.423761e-01	-1.946227e+00
25%	-4.115666e-01	-1.353324e-01	-2.362917e-02	-6.423761e-01	-6.582238e-01
50%	-4.115666e-01	-1.353324e-01	-2.362917e-02	-6.423761e-01	4.457787e-01
75%	-4.115666e-01	-1.353324e-01	-2.362917e-02	3.406008e-01	8.137795e-01
max	4.334796e+00	1.508160e+01	5.643531e+01	4.272508e+00	9.977799e-01

	...	52	53	54	55 \
count	...	3.823000e+03	3.823000e+03	3.823000e+03	3.823000e+03
mean	...	2.701940e-16	-9.007821e-16	3.516983e-16	4.302078e-16
std	...	1.000131e+00	1.000131e+00	1.000131e+00	1.000131e+00
min	...	-7.639066e-01	-1.932010e-01	-1.617539e-02	-3.050075e-01
25%	...	-7.639066e-01	-1.932010e-01	-1.617539e-02	-3.050075e-01
50%	...	-5.598673e-01	-1.932010e-01	-1.617539e-02	-3.050075e-01
75%	...	6.643687e-01	-1.932010e-01	-1.617539e-02	-3.050075e-01
max	...	2.500723e+00	1.543870e+01	6.182233e+01	1.047174e+01

	56	57	58	59	60 \
count	3.823000e+03	3.823000e+03	3.823000e+03	3.823000e+03	3.823000e+03
mean	4.029386e-17	4.445103e-16	3.178206e-16	-1.305086e-16	3.680318e-16
std	1.000131e+00	1.000131e+00	1.000131e+00	1.000131e+00	1.000131e+00

```

min    -1.176029e+00 -2.755685e+00 -2.296235e+00 -1.160247e+00 -5.227936e-01
25%    -9.752000e-01 -4.483164e-01 -4.930910e-01 -1.160247e+00 -5.227936e-01
50%    -1.718840e-01  2.438943e-01  3.083064e-01 -1.212969e-01 -5.227936e-01
75%     8.322610e-01  7.053680e-01  9.093544e-01  9.176534e-01 -2.623710e-02
max     2.037235e+00  9.361049e-01  9.093544e-01  1.610287e+00  3.449659e+00

```

```

count    61
count    3.823000e+03
mean    -1.956307e-15
std      1.000131e+00
min     -1.757406e-01
25%     -1.757406e-01
50%     -1.757406e-01
75%     -1.757406e-01
max      1.373073e+01

```

[8 rows x 62 columns]

```

[10]: # e
from sklearn.decomposition import PCA

# run the PCA algorithm on the standardized data with number of components
↳ equal
# to the total number of columns in the data.
pca = PCA(n_components=standarized_df.shape[1])
pca.fit(standarized_df)

```

[10]: PCA(n\_components=62)

```

[12]: # Fit the PCA model and transform data to get the principal components

# Instantiate PCA estimator
pca = decomposition.PCA()
df_plot = pd.DataFrame(pca.fit_transform(standarized_df), columns=
↳ standarized_df.columns, index=standarized_df.index)
df_plot

```

```

[12]:
      0      1      2      3      4      5      6  \
0    0.021179 -1.506218  4.028060  2.837064  1.121506 -1.048719  0.150893
1   -0.436318 -3.001971  6.068029  2.907716  1.439000 -0.447028  0.798956
2    1.363008  3.160016 -0.743226  1.395725  0.314563  0.968713 -3.625171
3    4.499442  0.949555  0.433865 -1.720058 -0.517172 -2.701682  1.721142
4   -1.199084 -3.264752  1.706263  1.130340 -1.262345 -0.225203 -1.368362
...    ...    ...    ...    ...    ...    ...
3818 -0.664700  0.199198  4.083232 -1.080674 -1.692163  1.007006 -0.950205
3819  4.295936 -4.344692 -2.048086  1.937549  3.064987  0.580706  0.111593
3820 -0.360348 -3.822805 -1.171858  1.853227 -0.492848  0.137883 -1.624683

```



```

3821 -1.351938 -5.283408 1.079286 3.439551 -1.187103 0.748811 -1.533931
3822 3.375207 4.220092 0.595156 0.048416 -1.199746 1.398314 -1.116004

```

```

      7      8      9  ...      52      53      54  \
0  -0.236166 -1.144235 0.216062 ... 0.079098 -0.099608 -0.328052
1  -1.762127 -1.696441 -0.287052 ... 0.595869 0.203963 0.451097
2  -2.136621 -0.731816 0.292576 ... -0.539342 0.126740 0.011906
3   0.417025 -1.923872 1.028773 ... 0.320966 0.120637 0.695168
4   2.425412 0.360641 -0.039846 ... -0.392310 -0.106501 0.015097
...
3818 0.498456 -0.628901 1.043451 ... 0.065532 0.205556 -0.269823
3819 -0.103625 -1.386905 -0.574874 ... -0.700666 -0.479226 -0.397264
3820 1.360577 1.415196 -1.159309 ... -0.172458 0.365674 0.195132
3821 1.703575 -0.440460 0.615567 ... 0.293833 -0.101132 -0.168334
3822 -1.452023 -1.034071 0.813605 ... -0.098148 0.095915 0.272171

      55      56      57      58      59      60      61
0  -0.049694 0.407308 -0.326841 0.361029 0.414270 0.096211 0.148911
1  -0.031893 -0.077154 0.484712 -0.462682 0.138005 -0.183608 -0.015544
2   0.141793 -0.578015 0.079724 0.067017 0.179884 -0.199439 -0.092629
3  -0.101969 -0.594626 0.927018 -0.535342 0.559234 -0.129050 -0.221793
4   0.002365 -0.544356 0.031109 -0.245522 -0.021291 -0.064236 -0.103955
...
3818 0.040811 0.132728 -0.035091 0.146050 -0.141946 0.395055 -0.486348
3819 0.298726 0.091618 0.063463 -0.079364 0.217436 -0.092169 0.106547
3820 -0.069563 0.207445 -0.378826 -0.302459 0.308900 -0.004402 -0.259517
3821 0.019024 0.007306 -0.090449 0.234468 0.101860 0.047271 -0.221078
3822 0.344237 -0.239525 0.327969 0.258709 -0.170725 0.013201 -0.359563

```

[3823 rows x 62 columns]

```
[13]: print(df_plot.shape)
      df_plot.iloc[:, 0]
```

(3823, 62)

```

[13]: 0      0.021179
      1     -0.436318
      2      1.363008
      3      4.499442
      4     -1.199084
      ...
3818  -0.664700
3819   4.295936
3820  -0.360348
3821  -1.351938
3822   3.375207

```

Name: 0, Length: 3823, dtype: float64

## 1 E

Let's look at the proportion of variance explained (PVE) of the raw data as explained by Principal Component Loading.

```
[14]: pca.explained_variance_ratio_
```

```
[14]: array([0.11639052, 0.10515794, 0.07626379, 0.05711962, 0.0496873 ,
          0.04648104, 0.03854584, 0.03379051, 0.02851031, 0.02630916,
          0.02503596, 0.02408691, 0.02278827, 0.02093577, 0.01955806,
          0.01890912, 0.01843386, 0.01800929, 0.01622157, 0.01565134,
          0.01425234, 0.01378334, 0.01288684, 0.01128819, 0.01061524,
          0.01012621, 0.00972703, 0.00929472, 0.00863325, 0.00824354,
          0.00787725, 0.00712658, 0.00695497, 0.00629398, 0.00607752,
          0.00592961, 0.00533857, 0.00488499, 0.00451338, 0.00446158,
          0.0043551 , 0.00401866, 0.00385701, 0.0034535 , 0.00332456,
          0.00309779, 0.00307227, 0.00277881, 0.00266685, 0.00256959,
          0.00242815, 0.00235895, 0.0021833 , 0.00207732, 0.00192919,
          0.0018067 , 0.0016325 , 0.00153693, 0.00142055, 0.00123766,
          0.00106562, 0.00093372])
```

From the above results, we see that the first principal component explains 11.6 % of the variance in the data, and the next principal component explains 10.5% of the variance and the third only 7 %. The variance explains tapers off as we move down the last PCA.


Even together, the first five principal components don't explain 50 % of the variance in the data.

So the PCAs will not be able to explain the variance in the data standalone. It can be depicted by the individual vs cumulative Proportion of Variance Explained in scree plot below.

```
[15]: # e ii. plot the Proportion of Variance Explained (PVE) of each principal_
      ↪ component (i.e., a scree plot)
      # and the cumulative PVE of each principal component.
      import matplotlib.pyplot as plt
      import numpy as np

      pve = pca.explained_variance_ratio_
      cumulative_pve = np.cumsum(pve)

      fig, ax = plt.subplots(figsize=(8, 6))
      ax.plot(range(1, len(pve) + 1), pve, label='PVE')
      ax.plot(range(1, len(pve) + 1), cumulative_pve, label='Cumulative PVE')
      ax.set_xlabel('Principal Component')
      ax.set_ylabel('Proportion of Variance Explained')
      ax.legend()
      plt.show()
```



```
[16]: # merging the first two pcas with the target variable.
pca_1_2_target = pd.concat([df_plot.iloc[:, 0], df_plot.iloc[:, 1],
    ↪ hand_digits[64]], axis=1)
pca_1_2_target = pca_1_2_target.rename(columns={0: 'pca_1', 1: 'pca_2', 64:
    ↪ 'handwritten_digit'})
pca_1_2_target.head()
```

```
[16]:      pca_1      pca_2  handwritten_digit
0  0.021179 -1.506218                0
1 -0.436318 -3.001971                0
2  1.363008  3.160016                7
3  4.499442  0.949555                4
4 -1.199084 -3.264752                6
```

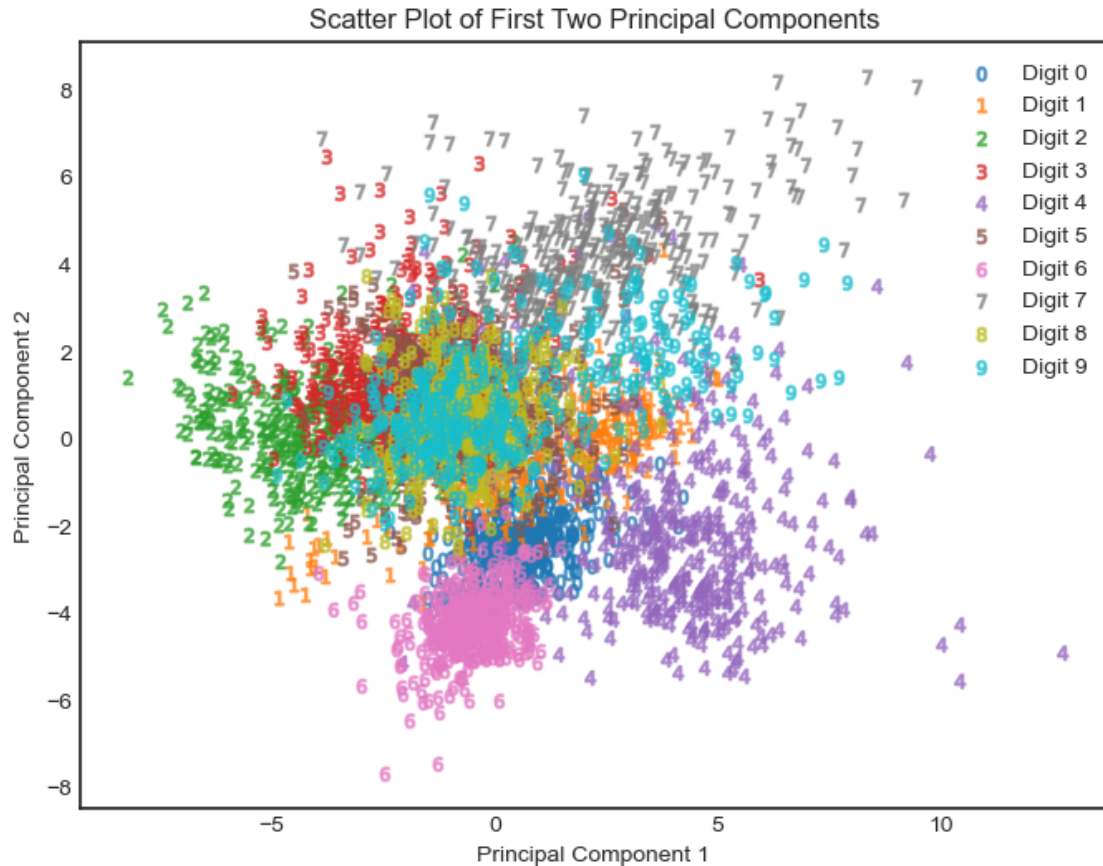
```
[17]: # scatter plot for the first two principal components
# the target class variable (i.e., digit number) with different symbols and
    ↪ colors
fig, ax = plt.subplots(figsize=(8, 6))
for digit in range(10):
    subset = pca_1_2_target[pca_1_2_target['handwritten_digit'] == digit]
```



```

ax.scatter(subset['pca_1'], subset['pca_2'], label=f'Digit {digit}',
           alpha=0.6, marker='${}$'.format(digit))
ax.legend()
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('Scatter Plot of First Two Principal Components')
plt.legend()
plt.show()

```



The target class variable, which represents the digit number, is represented by a different symbol– number itself– and color in the scatter plot.

The scatter plot shows how the data points are distributed in the two-dimensional space defined by PC1 and PC2. We observed distinct clusters for some of the digits like 6, 2, and 4. But the clusters are not well defined and bounded. The cluster for digit 6 is better explained by PCA1 while 3 and 2 are better explained by PCA 2.

It can also be due to reason that the two PCAs combine explain less than 25 % variance in the data. So we are not that confident in the effectiveness of the PCA in capturing the underlying structure of the data.